TDT4113 (PLAB2) Project 5: Keypad

— Assignment Sheet —

Purpose:

- Learn to properly interpret signals from a simple keypad.

- Learn about Charlieplexed circuits.

- Learn the basics of a Finite State Machine (FSM) and its use as a controller.

- Gain experience integrating the keypad, Charlieplexer and FSM on the Raspberry Pi (R-Pi).

Practical Information

- This project will be done in **groups of size three**.

- Each group will be issued a Raspberry Pi, keypad and related accessories for the TWO weeks of the project; these will be **returned to us immediately after the demonstration**.

- The due date for the demonstration is
  Wednesday, October 16, 2019 (kl 10:00)
  — we strongly recommend to prepare for presenting your solutions significantly earlier.

- Your code must be uploaded to BLACKBOARD prior to 8:00 (in the morning) on October 16.

- Your code needs to have a pylint level of at least 8.0

1 Introduction

As shown in Figure 1, your system, the *Keypad Controller*, will contain 3 key functional components: a finite state machine (FSM), the keypad and a *'Charlieplexed'* LED board.


Figure 1: Basic functional architecture of the keypad controller.

There is one central software object, the *agent*, which coordinates activity among them. The agent uses the keypad for input from the (human) user, and the LED board for output communication with the user. Within the agent, the FSM controls much of its decision making. However, as described later in this document, the FSM will work at a reasonably abstract level, leaving many of the straightforward, but detailed, computations (and their related data structures) to the agent.

You will implement the FSM as a simple rule-based system (RBS). Each behavioral rule consists of a condition and a consequent. During rule application, the condition is compared to the current situation (a.k.a. **context**) of the FSM. The FSM context consists of its internal state and its current signal. When *context* matches *condition*, the rule *fires* by performing the operations specified in the

consequent. These operations are simply changes to the FSM's state and requests for the agent to perform particular actions, such as verifying a password, lighting a particular LED, etc.

The keypad device does not come fully programmed to transfer simple symbolic signals (i.e. keystrokes) to your Python program, running on the Raspberry-Pi (R-Pi). Instead, you will need to explicitly probe the keypad's button array and discern digits (0-9) and two special symbols ("#" and "*") from the results of those probes.

1.1 Charlieplexing

As an added bonus, you will learn a creative trick for getting the most out of a breadboard that has limited number of connection pins to the Raspberry-Pi. The technique, known as *Charlieplexing*, enables you to control many actuators (e.g. lights, buzzers, etc.) using only a few input/output ports. The trick lies in the combination of wiring and Python programming.

1.2 Outline of the operation

The general operation of the agent is quite straightforward, with a complexity level similar to that of an advanced garage-door opener: a password-protected system enabling a few basic actions, including changing the password.

Via the keypad, the user will be able to enter sequences of symbols that will direct the agent to perform the following operations:

- Receive and validate — or reject — the password.
- Change the password.
- Select one of the 6 LEDs on the breadboard and turn it on for a time duration specified by the user via the keypad.
- Log out of the system.

We will now delve into the details of the FSM before discussing its interactions with the agent and the design of the keypad and Charlieplexed breadboard.

2 The Finite State Machine (FSM)

FSMs are often depicted as graphs: diagrams containing nodes with labeled, directed connections (a.k.a. arcs) between them. Each node represents a unique state of the system, while arc labels denote the external signal that triggers the transition from one state to another. As a simple example, Figure 2 provides two similar FSMs for identifying a particular password: 836. Presumably, the external signals, the password digits, come from a keypad.

The FSM begins in state S0. If the user presses an "8", the FSM transitions into state S1. It then moves to state S2 on a "3", and into the final state (S3) on a "6". Once in a final state, the FSM could then ask the agent to perform a password-protected action, such as unlocking a door, turning on a coffee machine, etc.

But what if the user makes a mistake? The upper FSM of Figure 2 specifies no particular actions for incorrect digits, which often implies that the FSM remains in the current state. But that is obviously a flawed design for a password reader, since the user could just keep guessing digits until she hits the proper one, thus advancing the FSM one state, and eventually pushing it to state S3.

Figure 2: Two simple Finite State Machines (FSMs) for detecting the password 836, with gray start state (S0) and red final state (S3). Symbol (Q) in box on arrow from state X to Y implies: *When in state X and receiving an external input (i.e. , from the keypad) of Q, transition to state Y.* (A) Upper FSM specifies state changes for only the correct password. (B) Lower FSM includes the "@" transition, which indicates "anything but the correct digit". For example, from state S2, the "@" indicates a transition back to S0 on anything but a 6.

A more reasonable design (Figure 2B), includes "@" transitions back to the start state, S0. These apply to any situation in which the user enters an incorrect digit; they force the user to start over. This is still somewhat *forgiving* in that it allows the user an unlimited number of attempts, but for this assignment, that level of security (or lack thereof) is perfectly sufficient.

Many alternatives for programming an FSM exist, including those in which each node of the graph is a node object, with arc objects emanating from it. Conversely, this assignment uses a **rule-based system**'s (RBS) approach to modeling the FSM. In an RBS, a set of rules (often called "IF-THEN" rules) are applied to the current context (i.e. FSM state plus trigger signal) to determine the new FSM state and (optionally) the action(s) to be performed by the agent.

An IF-THEN rule has this basic structure: "IF condition THEN consequent": the consequent is performed only if the condition (a.k.a. *antecedent*) is true. In the parlance of RBS, a rule is *applied* when its condition is compared to the current context, while a rule is *fired* when its condition matches the context and its consequent is performed. When a sequence of such rules are applied to a context, the standard assumption is that only ONE of them will fire, which will change the context, and then initiate another round of rule application. Furthermore, the **only** rule that fires is the **first** applied rule whose condition matches the context. Hence, the **ordering** of rules within the RBS is an essential factor that strongly influences its performance.

Although we will **not** implement our RBS as a big IF-THEN-ELSE statement, it helps to use such a statement in order to understand the logic behind an RBS. Using the FSM of Figure 2B as our model, the pseudocode below realizes that FSM's overall behavior:

> IF fsm.state == S0 AND fsm.signal == 8 THEN fsm.state ← S1
>
> ELIF fsm.state == S0 AND fsm.signal != 8 THEN fsm.state ← S0
>
> ELIF fsm.state == S1 AND fsm.signal == 3 THEN fsm.state ← S2
>
> ELIF fsm.state == S1 AND fsm.signal != 3 THEN fsm.state ← S0

ELIF fsm.state == S2 AND fsm.signal == 6 THEN fsm.state ← S3

ELIF fsm.state == S2 AND fsm.signal != 6 THEN fsm.state ← S0

Although this is pseudocode, the intention is that EXACTLY ONE of the rules will be applied, and it will be the FIRST rule whose antecedent matches the current context. This code can be simplified in several ways. First, if we assume that the forward-transition rule always appears before the "return to S0" rule in the list, then the antecedent of the latter rule can be simplified (i.e. generalized) as follows:

IF fsm.state == S0 AND fsm.signal == 8 THEN fsm.state ← S1

ELIF fsm.state == S0 THEN fsm.state ← S0

ELIF fsm.state == S1 AND fsm.signal == 3 THEN fsm.state ← S2

ELIF fsm.state == S1 THEN fsm.state ← S0

ELIF fsm.state == S2 AND fsm.signal == 6 THEN fsm.state ← S3

ELIF fsm.state == S2 THEN fsm.state ← S0

The important principle here is *specific-before-general*: the specific condition (e.g. fsm.state == S0 AND fsm.signal == 8) appears before the more general condition (fsm.state == S0). This insures that when the agent is in a particular state, the general rule fires only if the specific rule does NOT fire. Thus, the general rule effectively realizes the "everything but" logic signified by the "@" in Figure 2B.

Note that we could simplify the above pseudocode even more:

IF fsm.state == S0 AND fsm.signal == 8 THEN fsm.state ← S1

ELIF fsm.state == S1 AND fsm.signal == 3 THEN fsm.state ← S2

ELIF fsm.state == S2 AND fsm.signal == 6 THEN fsm.state ← S3

ELSE fsm.state ← S0

In this case, an equivalent RBS would house only 3 rules. Then, if none of them applied to the current context, the FSM's state would be reset to S0. Although this shortcut works fine, we normally want the contents of an RBS to mirror that of an FSM, and this is easiest if each arc of the FSM maps to a unique rule of the RBS. Hence, in Figure 2B, each "@" move should also have its own rule in the RBS. This ease of understanding becomes very important when large FSMs are to be implemented by an RBS.

2.1   Interactions between the Agent and FSM

As mentioned earlier, the FSM for this project will work at an abstract level, while the agent will do a lot of the routine computational work. Implementing all such activity within the FSM can be tedious, if not downright impossible. Essentially, we will rely on the agent to provide the FSM with a memory for everything other than the FSM's current state and signal.

We can begin with a simple practical aspect: a keypad controller system should allow the user to define and re-define their own password. For the time being, we will ignore exactly HOW the password gets defined but focus on how the FSM should be designed to check for ANY password, not just 836.

Assume that the agent has the current password (CP) stored in memory (or in a short text file). A more abstract FSM could then operate as follows:

- The FSM begins in state S-Init

- When the user hits any keypad button, the FSM switches to a "read password" state (S-Read).

- Each digit entered during this S-Read is simply passed to the agent, for storage in a cumulative-password (CUMP) array or string.

- When the user enters a '*' (asterisk), this tells the FSM to instruct the agent to verify the password, which it does by comparing CUMP to CP. The FSM also moves into a new state (S-Verify).

- The next signal from the agent should then indicate whether verification succeeded or not. Success causes the FSM to change to state S-Active, while failure forces a return to state S-Init.

The exact details of this sequence may vary, but the key point is that the FSM is no longer bound to a single password but can handle any password as long as the user signals the end of the password-entry phase via "*".

To handle agent-FSM interactions, the *consequent* aspect of each FSM rule needs to be expanded to a pair: new state, action. This means that the labels on arcs in FSM diagrams must change, but to a different pair: trigger symbol, action. Thus, the FSM updates its state and sends an action command to the agent. Figure 3 incorporates this enhancement in providing the FSM for a more flexible agent.
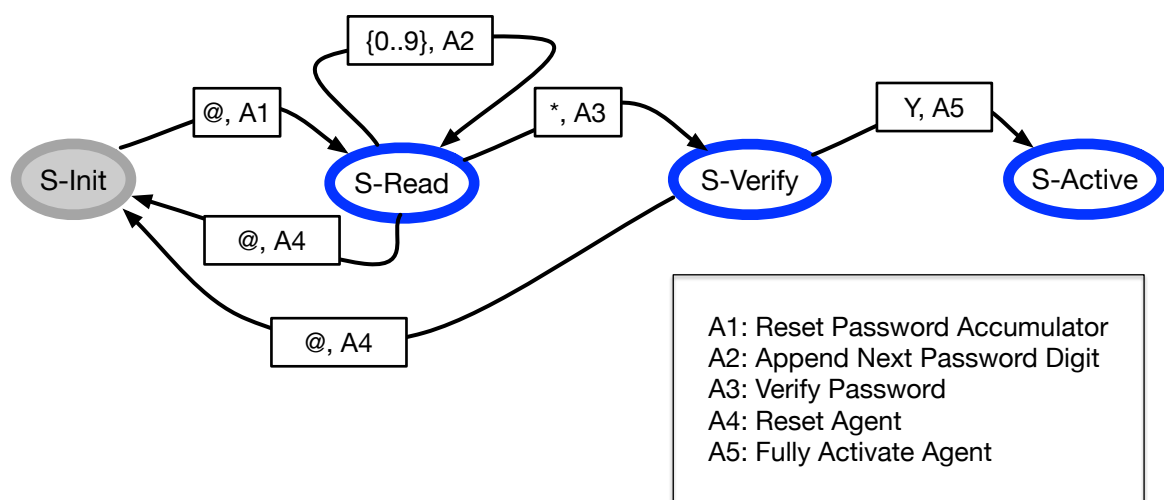

Figure 3: Portion of an FSM for a more abstract, and flexible, agent. Each arc indicates a pair: triggering symbol, agent action, with actions described in lower-right box. The symbol "@" still denotes "anything not specified as a trigger for the other outgoing arcs from this same state". For example, from state S-Verify, all symbols except "Y" (for yes/acceptance) send the FSM back to the start state, while **any** keystroke causes a transition from S-init to S-Read.

To extend this example to an FSM that accepts password changes, a few new FSM states and agent functions are required. Figure 4 illustrates this new FSM, which uses state S-active from Figure 3 as its starting point: once the user has successfully logged in via their (old) password, they are then in an FSM state (S-active) where it is possible to do many things, including changing to a new password.

The agent presumably has a few simple data structures for storing the current password (CP), the cumulative password (CUMP), and the last value of CUMP (CUMP-OLD). As shown in Figure 4, the user signals a desire to change the password by typing "*" (asterisk), which causes the FSM to change to state S-READ-2 and to instruct the agent to perform action A1, i.e. to reset CUMP. Each new digit that the FSM reads is then added to CUMP, via action A2. Upon reading a "*" in S-READ-2, the FSM initiates action A7, which will cache the current state of CUMP in CUMP-OLD and then reset/clear CUMP.

The FSM also transitions to state S-READ-3, whose behavior is very similar to S-READ-2, with the only difference being that upon receiving a "*" in S-READ-3, the FSM induces action A8, which will compare CUMP to CUMP-OLD and change the current password (CP) if and only if CUMP == CUMP-OLD. Either way, once A8 is done, the FSM returns to state S-Active. If the password does change, then the agent will need to write the new password to more permanent storage, such as a file.

In this example, the difference between agent actions R1 (reset agent) and R6 (refresh agent) may be quite minor. Resetting essentially means putting the agent in a very restricted mode where it can only receive password login attempts, whereas refreshing means returning to a less-restricted active state from which many other actions (not shown in Figure 4) such as turning on lights, may be initiated.
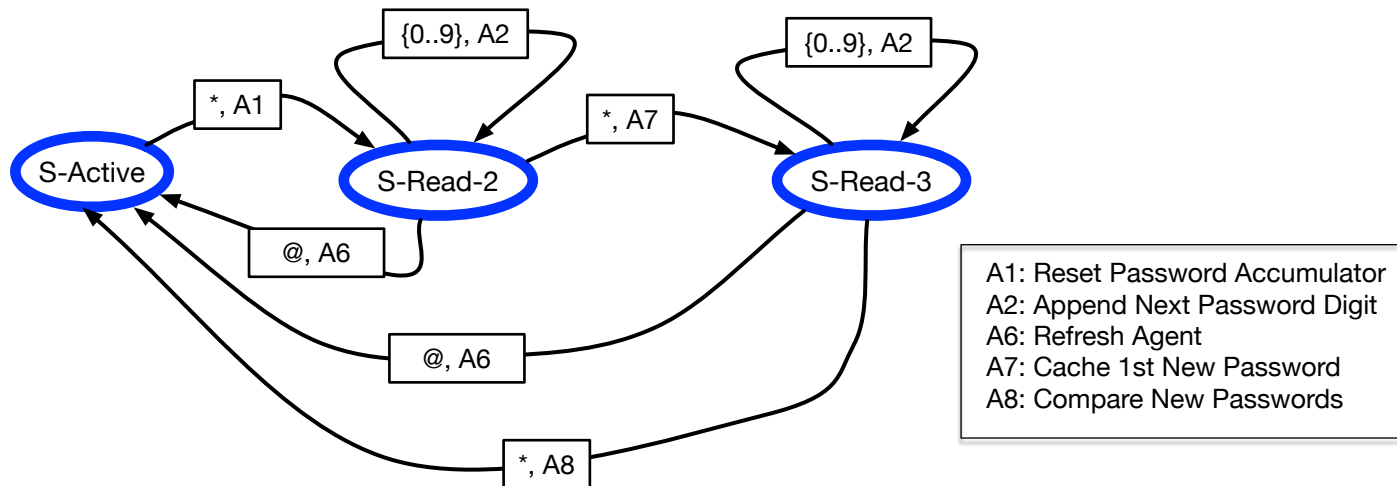

Figure 4: An extension of the FSM in Figure 3 to handle password updates, wherein the user must enter the new password twice.

2.2   Implementation Details of the FSM

To realize the FSM as a rule-based system (RBS), you will need to implement each rule as an object and to then cycle through those rules to determine the next action of the system.

Remember that a rule represents one arc (from an FSM graph) along with the nodes at its head and tail.

Thus, a rule object should contain at least these four instance variables:

1. state1 - triggering state of the FSM

2. state2 - new state of the FSM if this rule fires

3. signal - triggering signal

4. action - the agent will be instructed to perform this action if this rule fires.

Each rule has the following meaning:

> IF the FSM is in a state that matches rule.state1 and the current signal matches rule.signal THEN change the FSM's state to rule.state2 and call the method specified by rule.action (and give that method two arguments (the agent itself and the current signal).

In this description, note the word *matches*. This is not necessarily the same as *equals* but is a more general term.

Thus, rule.state1 may be a simple symbol, such as 's4', meaning that the FSM's actual state (S) must be s4. However, rule.state1 can also be a function that implicitly represents a collection of FSM states, any one of which might *match* S. Crucially, if rule.state1 is not a symbol, but a function, then that function needs to take exactly one argument: the FSM's actual state (S). The function then returns a boolean value (True/False) depending upon whether or not S matches its criteria. If the function accepts ANY state, then it should always return True, regardless of the value of S.

Similarly, rule.symbol can *match* the current symbol/signal (X) by either being a symbol that is identical to X, or by being a function that accepts X as an argument, performs some test on X, and returns a boolean. Or, if rule.symbol is meant to match anything, then it would simply accept X as an argument and output *True* regardless of X's value.

In another case, you may want rule.symbol to be a simple function that only accepts the digits 0-9. To do this in Python:

- Define a general function (not a method), such as:

    def signal_is_digit(signal): return 48 <= ord(signal) <= 57

- When defining the rule, simply write:

    rule.symbol = signal_is_digit

- Then, when checking whether the current symbol (sym) matches rule.symbol, make the call:

    rule.symbol(sym)

  which should only return True if sym is a digit.

Note, however, that rule.state2 should never be a function; it should always be a symbol specifying a single, legal state of the FSM. Similarly, rule.action should be a specific method that your agent object understands.

Practically speaking, you are probably much more likely to use a function for rule.symbol than for rule.state1. Still, you may have an FSM where the same transition can be taken FROM several different states, as exemplified by the "@" transitions in Figures 3 and 4; and in those cases, you may want to reduce the size of your rule base by having one rule (with rule.state1 as a function) that replaces several rules (with rule.state1 as different symbols in each rule). As mentioned above, this shortcut does destroy the one-to-one correspondence between the arcs of an FSM graph and the rules of an RBS, but you are free to explore either option.

Since rule.state1 and rule.symbol may be symbols or functions, at some point in your code you will probably want to import the *isfunction* function provided in Python's *inspect* module via this command:

    from inspect import isfunction

This allows you to quickly check whether rule.symbol (or rule.state1) is a function and to take action accordingly.

The following pseudocode reflects the basic operation of your overall system, with focus on the FSM:

- Create all rule objects for the FSM - and insure they are listed in the desired order.

- Initialize agent, keypad and LED board

- state ← fsm-start-state

- While state <> fsm-end-state do:

    symbol ← agent.get-next-symbol() ;; Get symbol, typically from the keypad

    For rule in fsm.rules: ;; (Rule-Application Loop)

    IF rule.match(state,symbol):

> state ← rule.state2
>
> agent.do-action(rule.action,symbol)
>
> GO TO (**)

>> End For
>
>> (**)

- End While

- Shutdown agent, keypad, LED board, etc.

As one useful option for coding up agent.do-action():

- Assume that your agent is of type KPC (keypad controller) and that KPC has several methods, including *start_password_entry*.

- Assume rule R3 declares: when in state S0 and reading symbol ″*″ (asterisk), transition to state S1 and invoke *agent.start_password_entry()*.

- R3 would then be defined as a rule with the following Python code used to set its action attribute:

> rule.action = KPC.start_password_entry

This syntax tells Python to fetch the desired method from the class KPC and put it into the action slot of rule.

- Given this rule definition, the actual Python syntax for agent.do-action() becomes:

> rule.action(agent,symbol)

where the first argument (agent) gets bound to *self* in the method's argument list.

- Then, the action slot of **every** rule in your FSM should be set to KPC.zzzz, where zzzz is a different method name.

Remember that the order of rules is extremely important. As a simple example, Table 1 lists the appropriate rules and their ordering for the FSM of Figure 3. In that table, note the use of KPC as the agent class and the use of simple functions such as *all_symbols(s)* that return True for any input symbol, s.

Table 1: A rule set corresponding to the FSM graph of Figure 3.

Index	State-1	Symbol	State-2	Action
1	S-init	all_symbols	S-Read	KPC.reset_password_accumulator
2	S-Read	all_digits	S-Read	KPC.append_next_password_digit
3	S-Read	*	S-Verify	KPC.verify_password
4	S-Read	all_symbols	S-init	KPC.reset_agent
5	S-Verify	Y	S-Active	KPC.fully_activate_agent
6	S-Verify	all_symbols	S-init	KPC.reset_agent

Note that rules 1, 4 and 6 use the *all_symbols* test, but that rule has a different meaning (a.k.a. semantics) in each case, due to the effects of rule ordering and the fact that we have declared specialized rules before the general rules. In rule 1, it really does mean ″all symbols″, but the effects of rule ordering are obvious when we get down to rule 4. Since rules 2 and 3 also apply in state S-Read, rule 4 will only be fired if rules 2 and 3 fail to match the current symbol. Hence, the actual semantics of *all_symbols* in rule 4 is ″any symbol except a digit or *″. So if the FSM is in state S-Read with a current symbol of ″*″, then rule 3 will fire, not rule 4, simply because the rule-ordering insures that rule 3 is always applied before rule 4....and only one rule can be fired for each round of rule application. Similarly, the semantics of *all_symbols* in rule 6 is ″all symbols except a Y″.

2.3   Additional Design and Debugging Tips

1. Although your FSM needs to be programmed as a rule-based system, it is **strongly advised** that you draw your complete FSM as a graph, and then convert the graph to a list of rules. It is much easier to debug your FSM logic in graph form than in rule form.

2. FSM states can be represented as simple strings or integers, while symbols/signals are probably easiest to model as strings (typically of length 1).

3. The method agent.get-next-symbol() should normally cause the agent to query the keypad (as detailed later in this document), but you should also keep an *override-signal* as a slot in the agent for cases where the agent's decision (such as to accept or reject a password) is the next signal that the FSM should receive, instead of keypad input. Be sure to clear the override-signal immediately after using it. The agent should know that if the override-signal is empty, it should query the keypad. Setting of the override-signal should only occur via your agent's action methods, and only a small subset of those methods should ever need to do so.

4. The call to agent.do-action() can be implemented in many different ways, but all legal agent actions should be defined as methods for your agent class.

5. Debugging of the FSM is much easier if done on your laptop with the physical keypad and LED board replaced by proxies. The proxy keypad can be a simple object that asks for user input from the laptop keyboard, while the proxy LED board can just print out messages such as "Turning light 3 on for 20 seconds".

3   Physical Setup of The Keypad Controller

As shown in Figure 5, several physical components contribute to the Keypad Controller. The keypad itself has many wires, each of which connects to the Transition Breadboard, whose only real purpose (when combined with the T-Cobbler) is to provide a more understandable interface between the keypad and the Raspberry Pi (as explained below).

All of the Python code for your system will reside on the Raspberry Pi. That includes code for your agent, FSM and interfaces to the keypad and LED board. The Charlieplexed LED Breadboard serves as the primary output for the agent. Different on/off patterns of the LEDs will provide basic (but useful) information to the human user.


Figure 5: Overall Physical Architecture of the Keypad Controller

4   The Keypad

As shown in Figure 6, the keypad consists of 12 keys arranged in a grid, with each key having a unique location (row, column). Since these keypads do not come with code for converting grid activity into keystrokes, you need to write that code yourself, as explained below. First, however, you need to make sure that the keypad is properly connected to your Raspberry Pi.
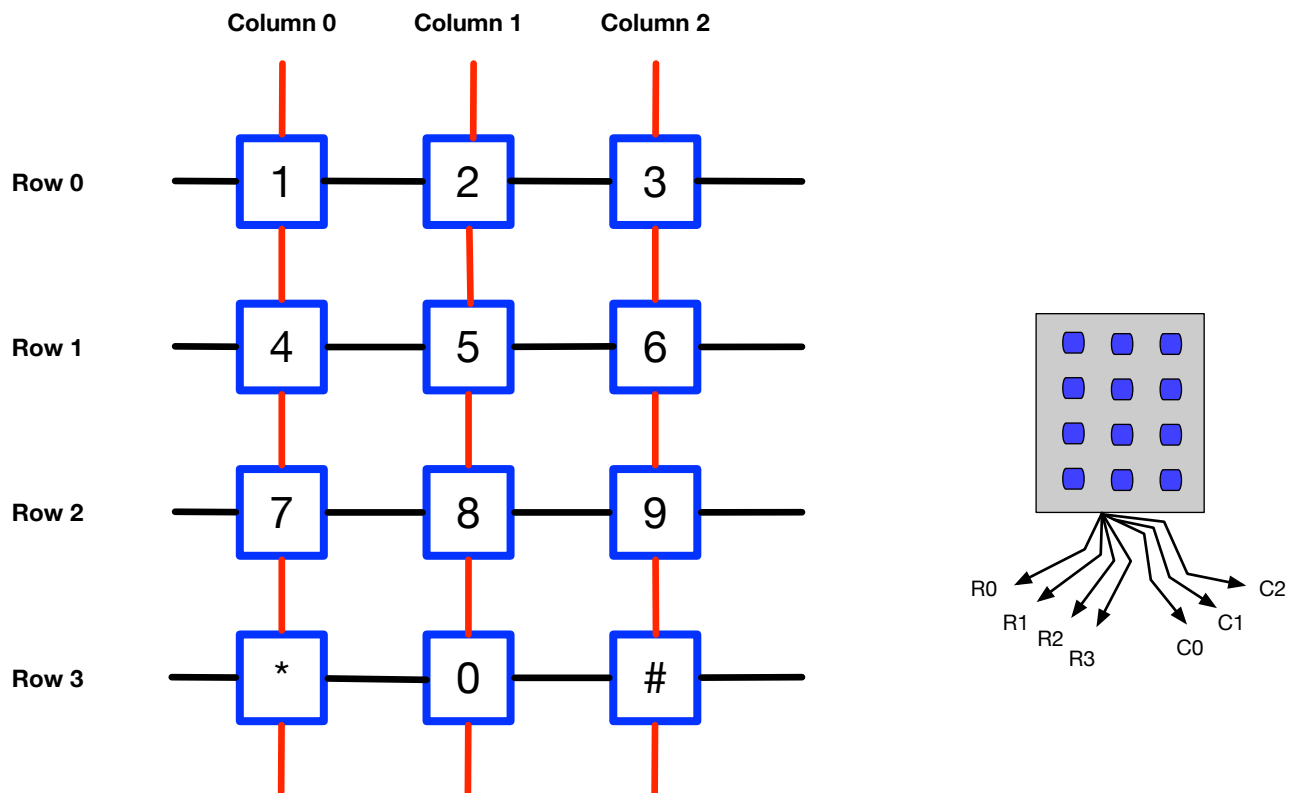

Figure 6: (Left) Internal circuitry of the keypad, with one wire for each row and column. (Right) External keypad circuitry, with all 7 wires exiting as a flat, transparent band from the bottom of the keypad. In that band, the ordering of row and column wires are as diagrammed: row 0's wire is leftmost, while column 2's is rightmost.

4.1   Connecting Keypad to Raspberry-Pi

For some odd reason, the pins on the Raspberry Pi (R-Pi) are not well labeled, along with being relatively difficult to access. For this assignment, we will use a T-Cobbler to essentially move our pin access points away from the crowded, unmarked R-Pi and out onto a spacious, well-marked breadboard. One end of the T-Cobbler matches up with the R-Pi pins, one-to-one, while the other end fits onto the transition breadboard AND provides legible pin labels. We can then connect the keypad directly to this breadboard and easily read off the pin numbers of the R-Pi that the keypad uses.

The keypad has one connection per row and column, so on our 12-key pad, we need $4 + 3 = 7$ pins into the R-Pi. For example, one pin configuration that seems to work well is: pins 18,23,24 and 25 for key rows 0,1,2, and 3, respectively; and pins 17, 27 and 22 for key columns 0, 1, and 2 respectively. See Figure 6 (right) for the ordering of the row and column wires as they exit the bottom of the keypad. These output wires are not labeled on the physical keypad.

4.2   Programming the Keypad

Your code should include a Keypad object that serves as an interface between the Keypad Controller agent and the physical keypad. The primary job of this object is to poll the keypad to detect keystrokes and to then send those keystrokes to the agent. The file for this code should also include the all-important declarations of which Raspberry Pi pins serve as inputs and outputs to the keypad code.

4.2.1   Pin Settings

Interactions between the R-Pi and keypad require the General Purpose Input/Output (GPIO) Python module for R-Pi. This should already be installed on your Raspberry Pi, so the import statement in your source code should look something like this:

> import RPi.GPIO as GPIO

The code snippets that follow assume that this import has been performed.

To initialize the row and column pins, these three operations are necessary:

1. GPIO.setmode(GPIO.BCM)

2. For each row pin, rp, do:

> GPIO.setup(rp,GPIO.OUT)

3. For each column pin, cp, do:

> GPIO.setup(cp,GPIO.IN,pull_up_down=GPIO.PUD_DOWN))

The use of PUD_DOWN declares that the pin will employ a pull-down resistor. This entails that a key push along the wire coming into that pin will cause a HIGH signal. The pull-down resistor seems to work better than the pull-up resistor for this model of keypad.

4.2.2   Polling the Keypad

To determine which (if any) key is currently being pressed, we will use a *polling* process in which different combinations of rows and columns are stimulated and sampled. This is implemented as a pair of nested loops, where the outer loop cycles through the keypad rows, setting them HIGH one at a time. Once a row (R) is stimulated, the inner loop cycles through the columns and reads them as input, one at a time. If column C has a HIGH reading, then the key at location (R,C) is revealed as the one being pressed. Simple calculations, and/or a dictionary can be used to map (r,c) pairs to the 12 symbols.

During the actual polling process, the command to set a row pin (rp) to high is:

> GPIO.output(rp,GPIO.HIGH)

and the code to check whether a column pin (cp) is high is:

> GPIO.input(cp) == GPIO.HIGH

Note that during polling, only one row pin should be high at a time. To reset a row pin to low, use:

> GPIO.output(rp,GPIO.LOW)

Finally, to avoid noisy inputs from the column pins, it helps to consider the column pin to be **actually high** only if repeated measurements (for example, 20 in a row with a 10 millisecond delay between each reading) all show a high value. You can use the time.sleep() command from Python's *time* package to support this simple (but very important) measure-wait-measure loop.

5 The Charlieplexed Breadboard

A Charlieplexed circuit board allows individual on/off control of many devices (e.g. LEDs) using just a few pins, where each pin has adjustable input/output settings. Consider the circuit diagram (without resistors) in Figure 7. If we declare pins 0 and 1 as output, and pin2 as input, then we are effectively allowing the voltages of pins 0 and 1 to drive the circuit, while pin2 nearly shuts down: not enough current flows through it to run devices along its portion of the circuit. This isolates board activity to LEDs 0 and 1.

Now, if we set pin0 high and pin1 low, then significant current will only flow through LED 0, lighting it.[1]. Conversely, if we set pin1 high and pin0 low, LED 1 lights up. Similar logic enables us to selectively light LEDs 2 and 3 by declaring pin 0 as input (neutralizing it) and pins 1 and 2 as output. Finally, with pin1 as input and pins 0 and 2 as output, LEDs 4 and 5 can be lit.

Charlieplexing permits the individual control of N(N-1) devices using only N pins. That relationship derives from a straightforward analysis: If you have N pins, then any pair of them can be used to control 2 different devices. There are "N choose 2" possible pairs of pins, which is:

$$\frac{N!}{(N-2)!(2)!} = \frac{N(N-1)}{2} \tag{1}$$

Multiply that by 2 (for 2 devices per pair), and you get N(N-1) total devices controllable by N pins. Thus, 4 pins are enough to control 12 LEDs, while 5 can control 20, etc. Of course, the circuit layouts for these larger configurations probably look like spaghetti; the nice planar layout of the 6-LED case cannot be maintained as N increases.


Figure 7: A Charlieplexed circuit with 3 pins and 6 LEDs. No resistors are drawn for ease of illustration. Arrows indicate the only legal direction of current flow for each LED, but these flows only occur under the proper pin settings; e.g. lighting LED 0 requires pin 0 to be high and pin 1 to be low.

5.1 Implementing the Charlieplexer

As with the keypad, the Charlieplexer should have an interface object that links it to the agent. Thus, all requests to perform different lighting sequences on the LED Board should go from the agent to the board via this interface object.

Once the 6-LED Charlieplexed circuit is wired up (see http://razzpisampler.oreilly.com/ch04.html for details), the Python coding is straightforward. Using RPi.GPIO, we can easily change pin modes back and forth between input and output. So whichever of the 6 LEDs we want to light,

---

[1]LEDs permit current flow in only one direction

there will be two pins chosen as output and one for input; and one of the output pins will be set HIGH, the other LOW. The pin settings corresponding to each LED should be saved in a dictionary or array such that when the agent requests lighting of the kth LED, your code can fetch the kth setting set and send it to a method that performs the proper in/out and high/low assignments. All of this can be done with 10 or 15 lines of Python code.

Unlike the keypad, the Charlieplexer for this assignment requires no special declaration of a pull-up or pull-down resistor when declaring a pin as an input. Thus, these basic commands work fine for declarations of input and output:

- GPIO.setup(pin,GPIO.IN)

- GPIO.setup(pin,GPIO.OUT)

To set the voltage level of an output pin (outpin), the same commands as used for the keypad work fine:

- GPIO.output(outpin,GPIO.HIGH)

- GPIO.output(outpin,GPIO.LOW)

The following are mandatory behaviors that your LED board needs to perform:

1. A display of lights (of your choosing) that indicates "powering up". This should be performed when the user does the very first keystroke of a session.

2. Flashing of all lights in synchrony when the user enters the wrong password during login.

3. Twinkling the lights (in any sequence you choose) when the user successfully logs in.

4. A fourth type of light display (of your choosing) that indicates "powering down". This should be performed immediately after the user logs out.

5. Turn one user-specified LED on for a user-specified number of seconds, where information about the particular LED and duration are entered via the keypad.

Each of these activities should be initiated by a different method, callable by the agent.

Optional: Add other LED behaviors (of your choosing), triggered by other keypad symbols, to this basic behavioral repertoire.

6   Primary Python Classes, Instance Variables, and Methods

The following section provides the basic class structure and functionality, but you will almost certainly implement more methods (and possibly more classes) as you see fit.

The four main classes should be:

1. Finite State Machine (FSM)

2. Keypad - interface to the physical keypad

3. Led Board - interface to the physical, Charlieplexed LED board.

4. KPC - the keypad controller agent that coordinates activity between the other 3 classes along with veryifying and changing passwords.

6.1   Finite State Machine

An FSM object should house a pointer back to the agent, since it will make many requests to the agent (KPC) object.

Key methods for the FSM include:

- add_rule - add a new rule to the end of the FSM's rule list.

- get_next_signal - query the agent for the next signal.

- run_rules - go through the rule set, in order, applying each rule until one of the rules is fired.

- apply_rule - check whether the conditions of a rule are met.

- fire_rule - use the consequent of a rule to a) set the next state of the FSM, and b) call the appropriate agent action method.

- main_loop - begin in the FSM's default initial state and then repeatedly call get_next_signal and run_rules until the FSM enters its default final state.

The actual arguments to these methods may vary depending upon your exact implementation.

6.2   Keypad

The keypad has a few essential methods:

- setup - Set the proper mode via: GPIO.setmode(GPIO.BCM). Also, use GPIO functions to set the row pins as outputs and the column pins as inputs.

- do_polling - Use nested loops (discussed above) to determine the key currently being pressed on the keypad.

- get_next_signal - This is the main interface between the agent and the keypad. It should initiate repeated calls to do_polling until a key press is detected.

6.3   LED Board

The Charlieplexed LED Board has these methods:

- setup - Set the proper mode via: GPIO.setmode(GPIO.BCM).

- light_led - Turn on one of the 6 LEDs by making the appropriate combination of input and output declarations, and then making the appropriate HIGH / LOW settings on the output pins.

- flash_all_leds - Flash all 6 LEDs on and off for k seconds, where k is an argument of the method.

- twinkle_all_leds - Turn all LEDs on and off in sequence for k seconds, where k is an argument of the method.

In addition, you will need methods for the lighting patterns associated with powering up (and down) the system.

6.4   KPC Agent

A KPC agent is the main object in your system, so it will have many methods and a few important instance variables. These variables include:

- a pointer to the keypad,

- a pointer to the LED Board,

- a few simple strings or arrays for holding important sequences of keystrokes, such as a passcode-buffer for all numbers in an ongoing password-entry attempt.

- the complete pathname to the file holding the KPC's password.

- the override-signal (discussed earlier in this document)

- slots for holding the LED id (Lid) and lighting duration (Ldur) – both entered via the keypad – so that it can initiate the action of turning a specific LED on for a specific length of time.

Some of the most important KPC methods are the following:

- init_passcode_entry - Clear the passcode-buffer and initiate a "power up" lighting sequence on the LED Board. This should be done when the user first presses the keypad.

- get_next_signal - Return the override-signal, if it is non-blank; otherwise query the keypad for the next pressed key.

- verify_login - Check that the password just entered via the keypad matches that in the password file. Store the result (Y or N) in the override-signal. Also, this should call the LED Board to initiate the appropriate lighting pattern for login success or failure.

- validate_passcode_change - Check that the new password is *legal*. If so, write the new password in the password file. A legal password should be at least 4 digits long and should contain no symbols other than the digits 0-9. As in verify_login, this should use the LED Board to signal success or failure in changing the password.[2]

- light_one_led - Using values stored in the Lid and Ldur slots, call the LED Board and request that LED # Lid be turned on for Ldur seconds.

- flash_leds - Call the LED Board and request the flashing of all LEDs.

- twinkle_leds - Call the LED Board and request the twinkling of all LEDs.

- exit_action - Call the LED Board to initiate the "power down" lighting sequence.

7 A Typical Run of the Keypad Controller

This section walks through one sample session with the Keypad Controller, showing what the user types and how the system responds.

The session begins with the Raspberry Pi being turned on and the main controller program started up in Python. Once started, the controller will wait for input from the agent. Table 2 shows a sequence of agent inputs (most of which will come from the keypad) and the corresponding LED-board actions and FSM states. These are the states that the FSM moves into after reading input from the agent. For inputs of length 2 or longer, it is assumed that the FSM repeats the corresponding state (such as S_{read} during password entry). Note that this is only an example: the FSM that you design may involve different states and connections between them.

In this example, the complete sequence of keypad inputs is:

8 12345 * 4 * 29 * 2 * 14 * # #

whereas one symbol "Y" (4th table entry) was provided as an override-signal from the agent, thus signaling password acceptance.

In Table 2, the FSM state S_{active} is the main start state for most activities, but successful login is required to reach S_{active}. Also note that all of the agent's inputs come from the keyboard **except** the "Y", which is the agent's *override signal* to the FSM that the password has been accepted. Similarly, an "N" should indicate password rejection and send the FSM back to state S_{init}. Finally, notice that logging out requires two pushes of the # key, one to move from state S_{active} to S_{logout}, and a second to confirm the desire to logout.

---

[2]For this assignment, it is sufficient to ask the user to type the new password just ONCE, although twice, as shown by the FSM in Figure 4, is more realistic.

Table 2: Example sequence of agent' inputs (primarily from the keypad) and system responses.

Activity	Input from Agent	Example	LED Display	FSM State
Wake up System	Any Key Press	8	Power-up light show	S_{init}
Enter Password	Digits of Password	12345	None	S_{read}
Completing Password Entry	*	*	None	S_{verify}
Password Accepted	Y	Y	Twinkling lights	S_{active}
Choose a LED	Digit in 0-5	4	None	S_{led}
Begin Duration Entry	*	*	None	S_{time}
Choose a Duration	Digits	29	None	S_{time}
Complete duration	*	*	LED 4 ON for 29 secs	S_{active}
Choose a LED	Digit in 0-5	2	None	S_{led}
Begin Duration Entry	*	*	None	S_{time}
Choose a Duration	Digits	14	None	S_{time}
Complete duration	*	*	LED 2 ON for 14 secs	S_{active}
Begin Logout	#	#	None	S_{logout}
Confirm Logout	#	#	Power-down light show	S_{done}

8 Demonstration of the Keypad Controller

During the demonstration session for this project, you will be asked to perform the following sequence of actions using the keypad:

1. Press any key and a "powering up" light display begins.

2. Attempt to login, but fail, thus triggering the "failed login" light display.

3. Successfully login, thus triggering the "successful login" light display.

4. Change the password - your system must accept all-digit passwords of length 4 or greater.

5. Logout, thus triggering the "powering down" light display.

6. Press any key to begin the "powering up" light display.

7. Login again, using the new password.

8. Turn on a user-specified LED for a user-specified number of seconds.

9. Turn on a different LED for a different duration.

10. Logout, thus triggering the "powering down" light display.

9 Appendix: Running KPC without a screen

This is an optional part of the assignment. You do not have to do it to get full credit for the project.

Many electronic devices, such as garage-door openers, use keypads but do not provide any form of textual interface: their signals (if any) to the user are very primitive. Your KPC is designed to work without a screen, since the primary communication channels between the human user and your system are the keypad and the LED board, whose varied light patterns convey information. However, you will normally use a screen to open a terminal window, move to the proper code directory, and start python.

Running screenless entails getting your KPC agent to start running as soon as your Raspberry Pi boots up. There are many ways to achieve this; below is one relatively straightforward solution consisting of several steps:

1. Add a special "main" declaration in the file for your KPC agent class (e.g. kpc.py)

2. In the file /etc/rc.local, add a call that starts up python3 and runs that main routine.

3. Insure that the full pathname to your password file is known to the KPC agent.

In the following description, assume that the directory holding all of your code for this project is named "p5dir" (project 5 directory), and it resides just below the default main directory, which is probably /home/pi, /home/robot, or /home/keypad (depending upon how the system staff set up your R-Pi).

To add a "main" declaration to your KPC code, go to the bottom of the file (kpc.py) and add this line:

if __name__ == "__main__": my_startup_function()

This insures that when you issue the following command in unix (linux):

python3 kpc.py

python3 will start up, load in kpc.py and all other imports declared in kpc.py, and then call your startup function, my_startup_function. To test that this is working properly, while your R-Pi is still connected to a screen, go to directory /home/pi/p5dir and enter that command: "python3 kpc.py". Your system should start up, and your keypad should be waiting for you to press a key.

Next, you need to edit the file /etc/rc.local so that it includes a similar unix command. To do so, go to directory /etc on your R-Pi and use an editor such as nano (you probably need to use sudo to get permission to access and modify the file):

sudo nano rc.local

Just before the final line of rc.local (which is "exit 0"), add this line:

sudo python3 /home/pi/p5dir/kpc.py &

Note that you need the full pathname to the kpc.py file, since it's dangerous to assume a particular home directory in the middle of the whole boot-up process. The ampersand (&) at the end tells linux to fork a process for your system while continuing with its normal boot-up operations.

Finally, since the startup call to your system will come from outside the main directory for your code (p5dir), the complete path to your passcode file (i.e. , the file that holds nothing but the KPC password; let's call it "passcode.txt") should be used by your KPC code. So even though passcode.txt will probably reside in directory p5dir, because your code is started from a different directory during boot-up, the system will need to know the full path to passcode.txt. The simple remedy is to make sure that all references to the passcode file in your code use the full pathname.

— o —